

PROGRAMMING ASSIGNMENT #5
CS 2223 B-TERM 2022
CLOSED HASHING
&
DIJKSTRA'S ALGORITHM

ONE HUNDRED POINTS
DUE: THURSDAY, DECEMBER 8, 2022 11 PM

1. (10 Points) Read a file and hash the words:
Read a text file (EdgarAllanPoeBellsB2022groomed.txt from Canvas) and hash the words using the suggested hash function from *Levitin*, page 269:

$$h \leftarrow 0; \text{ for } i \leftarrow 0 \text{ to } s - 1 \text{ do } h \leftarrow (h * C + \text{ord}(c_i)) \bmod m,$$

where:

h is the computed hash value,
 s is the length of the word being hashed,
 c_i is the i^{th} character of the word,
 $\text{ord}(c)$ is the numerical value of character c in the alphabet in use,
 C is a constant larger than every $\text{ord}(c_i)$, and
 m is the modulus defining the size of our hash table.

We will take $\text{ord}(c)$ to be the ASCII value of the characters in each word. Thus, $\text{ord}(c) \in \{39, 45, 65, 66, \dots, 90, 97, 98, \dots, 122\}$ for $c \in \{', -, A, B, \dots, Z, a, b, \dots, z\}$, respectively. All other characters are discarded; we define words as unbroken sequences (strings!) of consecutive alphabetic characters, both upper and lower case, plus apostrophes and hyphens¹, in the lines read for the file. (Java uses Unicode, but ASCII is still under there somewhere—you just need to dig it out.)

For our “EdgarAllanPoeBellsB2022groomed.txt” text file, we will take $C = 123$ and $m = 293$. These represent parameters that can be changed to alter the shape and composition of our resulting hash table. For now, we will use these values for the sake of consistency. You are free to experiment with changing these parameters on your own, of course!

¹Thus, *turtle-dove* and *don't* are each individual words, though the latter does not appear in the poem.

2. (20 Points) Create a Hash Table using Closed Hashing (Open Addressing):

- a. Build a hash table of size 293, i.e. entries from 000 to 292, from the hash values computed in Part 1 above. Load them into the table in the order they occur in the file, *discarding duplicates*. Our table will get close to full, but you needn't worry about re-sizing it—we will examine the effects of its load factor in Part 3 below.
- b. Display the table, starting with table entry 0, in lines of the form:

Hash Address, Hashed Word, Hash Value of Word

(Remember, the Hash Address will not match the Hash Value when the resolution of a collision forces a word to be cascaded down the table. Also, remember that the table should be thought of as a circular list so that a word which cascades past the bottom end of the table gets wrapped to the top in looking for an open place in the table.)

Note: Your program may perform Parts 1 & 2 simultaneously, if you wish.

3. (30 Points) Explore the Hash Table:

Add to your code routines that answer these questions. Some answers may not be unique. You may resolve these issues by presenting either *any* correct answer or *all* correct answers.

- a. How many non-empty addresses are there in the table? What does that make the load factor, α , for our table?
- b. What is the longest empty area in the table, and where is it?
- c. What is the longest (largest) cluster in the table, and where is it?
Note: It might wrap from the end of the table back to the beginning.
- d. What hash value results from the greatest number of distinct words, and how many words have that hash value?
- e. What word is placed in the table farthest from its actual hash value, and how far away is it from its actual hash value?

I'll provide sample answers from a couple of different text files for you to check against over the weekend.

4. (40 Points) Implement Dijkstra's Algorithm with weighted graphs like this:

Your code should read a text file² and then ask for input from the console for start and destination nodes and use Dijkstra's algorithm to find the "length" of the shortest path between them. It should also display the sequence of nodes that constitute this shortest path.

You can assume the associated graph, read from a text file, is connected and that the input file is consistent—the main diagonal will consist exclusively of zeroes, there will be no negative edges, etc. The "10" atop the matrix below is a sentinel signaling the size of the graph to be loaded.

| | | | | | | | | | |
|----|-----|----|----|-----|-----|-----|-----|-----|-----|
| 10 | | | | | | | | | |
| 0 | 53 | 10 | 12 | 0 | 0 | 0 | 0 | 0 | 0 |
| 53 | 0 | 33 | 0 | 2 | 0 | 101 | 0 | 0 | 0 |
| 10 | 33 | 0 | 9 | 30 | 18 | 0 | 0 | 0 | 0 |
| 12 | 0 | 9 | 0 | 0 | 17 | 0 | 0 | 6 | 0 |
| 0 | 2 | 30 | 0 | 0 | 14 | 123 | 122 | 0 | 0 |
| 0 | 0 | 18 | 17 | 14 | 0 | 0 | 137 | 7 | 0 |
| 0 | 101 | 0 | 0 | 123 | 0 | 0 | 8 | 0 | 71 |
| 0 | 0 | 0 | 0 | 122 | 137 | 8 | 0 | 145 | 66 |
| 0 | 0 | 0 | 6 | 0 | 7 | 0 | 145 | 0 | 212 |
| 0 | 0 | 0 | 0 | 0 | 0 | 71 | 66 | 212 | 0 |

²You can hardcode this matrix if you prefer—we have a lot to get done in a short time.

